

Institut für Informatik
Lehrstuhl für Programmierung und Softwaretechnik

LUDWIG-MAXIMILIANS-UNIVERSITÄT MÜNCHEN



Bachelor Thesis

MicroPsi and Minecraft

A Popular Video Game as a Simulation Environment
for a Cognitive Architecture

Jonas Kemper

Medieninformatik Bachelor

Aufgabensteller: Prof. Dr. Martin Wirsing
Betreuer: Joscha Bach PhD, Annabelle Klarl
Abgabetermin: 19. September 2013

Ich versichere hiermit eidesstattlich, dass ich die vorliegende Arbeit selbstständig angefertigt, alle Zitate als solche kenntlich gemacht sowie alle benutzten Quellen und Hilfsmittel angegeben habe.

München, den 19. September 2013

.....
(Unterschrift des Kandidaten)

Zusammenfassung

Simulationsumgebungen spielen für das Testen und Erforschen von neuen Ansätzen für künstliche Intelligenz eine entscheidende Rolle. Da sich Computerprogramme nur mit erhöhtem technischem und finanziellem Aufwand in die physische Welt implementieren lassen, können simulierte Umgebungen schnellere, günstigere und reproduzierbarere Ergebnisse liefern. Genauso wie Innovationen aus dem Bereich der KI neue Impulse setzen, können auch neue Ansätze für Simulationsumgebungen zu neuen Erkenntnissen führen.

Diese Arbeit präsentiert eine Schnittstelle, welche es der kognitiven Architektur MicroPsi 2 ermöglicht, Agenten in Umgebungen des populären Videospiels Minecraft zu steuern. Minecraft bietet sich als Grundlage für eine Simulationsumgebung aufgrund der kompositionalen Semantik der Spielwelt besonders an, da das Agentensystem, durch das Erforschen seiner Umwelt, Wissen über die Spielwelt aufbauen und ähnliche Strukturen wiedererkennen kann. Objekte der Spielwelt sind in Minecraft keine bloßen Hindernisse, sondern werden mit variablen Eigenschaften prozedural erzeugt und erzeugen hierdurch Welten, die einer realen Umgebung mehr ähneln als andere virtuelle Welten.

Da Minecraft von Anfang an mit einem Mehrspielermodus ausgestattet wurde, lässt es sich zudem für Multiagenten-Umgebungen und so für kolaborative Agenten verwenden. Daraüber hinaus sind Minecraft-Lizenzen günstig zu erwerben, erhältlich für viele Plattformen und es gibt eine äußerst große und aktive Community für selbsterstellte Spiel-Inhalte und -Modifikationen.

Die für diese Arbeit entwickelte Schnittstelle ermöglicht es MicroPsi, sich an einem Minecraft Server anzumelden, die Umgebung wahrzunehmen und sich fortzubewegen. Darauf aufbauend wurde eine Visualisierung der Umgebung aus Sicht des Agenten implementiert, welche als OpenGL gerenderte 3D-Ansicht live im Browser angezeigt wird. Zum Schluss der Arbeit wird ein einfaches Experiment durchgeführt, bei welchem sich der Agent selbstständig auf ein bestimmtes Objekt zubewegen muss. Dieses Experiment wurde als Teil der Arbeit umfangreich dokumentiert.

Abstract

Simulation environments play an important role for testing and researching new approaches to artificial intelligences. Because computer software can only be implemented into the physical world with increased technical and financial effort, simulated environments can deliver results faster, cheaper and more reproducible.

In the same way as innovations in AI can deliver new impulses, new approaches to simulation environments can just as well lead to new insights into the future of intelligent machines..

This thesis presents an interface, which enables the cognitive architecture MicroPsi 2, to control agents in environments of the popular video game Minecraft.

Because of its compositional semantic, Minecraft turns out to be an interesting simulation environment, as agents can generate knowledge through exploring their environment and may recognise similar structures. Objects in Minecraft worlds are not just obstacles, but are generated procedurally with varying characteristics and therefore make the worlds more similar to real environments than other virtual worlds.

Since Minecraft contained a multi player mode from its early days on, it is also suitable for multi agent environments and therefore for collaborative agents. Furthermore, Minecraft licenses are affordable and available for many platforms. Also, there is a huge and active community for player generated content and modifications.

The interface enables MicroPsi, to connect to a Minecraft server, perceive its environment and move around within it.

Additionally, a visualisation of the environment, as it is perceived by the agent, has been implemented. It displays an OpenGL rendered 3D view live in the web browser.

Finally, a simple experiment is concluded, in which the agent has to move towards a specified object. This experiment is documented as a part of this thesis.

Acknowledgements

I would like to thank Prof. Dr. Martin Wirsing for giving me the opportunity to prove my motivation in writing a thesis about what I considered would suit me best. I thank my supervisor Joscha Bach, for giving me the opportunity to work on his project, advising and inspiring me through the entire process. I especially thank my supervisor Annabelle Klarl, for supporting my plan without hesitation, guiding me through the process and providing valuable feedback whenever I needed it. Dominik Welland helped me with understanding and building upon the MicroPsi framework and fixed inhabitant bugs as soon as I reported them, thank you!

Additionally I would like to thank Michael Fogleman whose code I built upon and especially Nick Gamberini who allowed me to use his Spock for this project and who answered every question about it in detail. Finally, I thank all the other members of the Minecraft community, who collect and distribute structured knowledge regarding the game in a combined effort and have answered my questions about the game and how it works in IRCs and message boards.

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Outline	2
2	Artificial General Intelligence	3
2.1	Strong AI	3
2.2	Psi Theory	4
2.3	Agents and Environments	7
2.4	Psi Implementations	8
2.5	MicroPsi 2	9
2.5.1	Module Overview	9
2.5.2	Module Server	10
2.5.3	Module Runtime	10
2.5.3.1	Node Nets	10
2.5.3.2	Worlds	11
2.6	Summary	12
3	Minecraft	15
3.1	Overview	15
3.1.1	Main Ideas	15
3.1.2	Development Process	16
3.1.3	Reception	16
3.2	Game Mechanics	16
3.2.1	Minecraft Worlds	17
3.2.2	The Client Server Protocol	17
3.3	Minecraft as a Gaming Platform	19
3.4	Suitability of Minecraft as a Simulation Environment	19
3.5	Related Projects	20
3.5.1	Minecraft Bots	20
3.5.1.1	Spock by Nick Gamberini	20
3.5.1.2	Protocol Implementation in Spock	21
3.5.2	Minecraft Clones	21
3.6	Summary	22
4	Minecraft as a MicroPsi 2 World	23
4.1	Overview	23
4.2	Using Spock as the <code>minecraftClient</code>	24
4.2.1	Overview of the <code>minecraftClient</code>	25
4.2.2	Extensions to the <code>minecraftClient</code>	26

4.3	Module <code>minecraftVisualisation</code>	26
4.4	Module <code>minecraftWorld</code>	27
4.5	Case Study	28
4.5.1	Evaluation	29
4.6	Summary	30
5	Conclusion	33
	List of Figures	35
	List of Tables	37
	Contents of the CD	39
	Bibliography	41

Chapter 1

Introduction

The hunt for artificial intelligence (AI) started many years ago. Dividing the subject into the strong and the weak part means separating its goals into developing useful, working applications and learning about the nature of intelligence itself. The ultimate task of strong AI—recreating human intelligence—admittedly still seems to be science fiction, though.

But then again: a new generation of cognitive scientists, psychologists and computer scientists strives to implement new ideas for simulated cognition by building cognitive architectures. Many of them do so by simulating, in one way or the other, what we would call neural node nets.

One of these architectures is MicroPsi 2 [Bac12]. Developed by Joscha Bach and Dominik Welland, it is based on the Psi theory by Dietrich Dörner [?], who is a German professor for theoretical psychology. It aims at providing a reliable and complete implementation of his theory of cognition and at the same time being easily accessible, understandable and modifiable for the research and applications to come.

To test the functionality of such cognitive architectures and to figure out their capabilities and potential, we need to study their behaviour inside defined environments. As implementing AI into the physical world (as robots, for example) requires patience and building appropriate hardware, computer-simulated environments play an important role.

MicroPsi 2 is both: a cognitive architecture as well as a set of and an interface to simulation environments.

1.1 Motivation

Video games are natural applications of artificial intelligence. The quality of a game's AI can make the difference between an immersive experience and an uninspired demonstration of computer graphics technology. One game that especially stood out in the recent years is Minecraft. As a so called *sandbox game*, it attempts to leave the player with as few limitations as possible, for what they do and how they behave in the game world. This kind of relatively open world delivers a simulated environment rich of opportunities, for both human and artificial players. Building an interface in between the cognitive architecture MicroPsi 2 and the video game Minecraft is what this thesis is about.

1.2 Outline

This project is fundamentally about combining existing technologies. The most important ones being MicroPsi 2, the framework aiming to implement the ideas of the Psi theory, and Minecraft, the popular sandbox videogame.

To understand how and why they were chosen, a brief history of their evolution as well as explanations of their basic ideas and relevant insights into their architecture are given in chapter 2 and 3. Chapter two introduces Artificial General Intelligence and in particular the Psi theory and its implementations. We will especially focus on the particular implementation MicroPsi 2 and describe its modules in detail. Chapter 3 introduces the video game Minecraft, its protocol and related software projects. As the resulting software heavily relies on them, the protocol and the utilised open source software will get increased attention. In Chapter 4 the contribution of this thesis, the Minecraft world adapter for MicroPsi 2, is being described in detail. As a proof of concept, the description eventually leads to a small experiment. Chapter 5 summarises the findings of this thesis and gives an outlook on possible future applications.

Chapter 2

Artificial General Intelligence

To provide the necessary context for this thesis, this chapter will provide a brief introduction to the corresponding foundations on the AI side.

First, we will introduce AIG as a subset of artificial intelligence. Based on this, the Psi theory, as a particular instance of strong AI theories, is described in detail. The description includes a brief history of its concrete implementations. Then, MicroPsi 2, the latest framework dedicated to the Psi theory, is described in detail.

2.1 Strong AI

It took many years for AI research to evolve from the early ideas of thinking machines over Deep Blue, the computer that could beat mankind's best chess players to Watson, the AI that beats the champions of Jeopardy, the game show, that is about asking the appropriate question to a given answer. There exist many applications for AI. Self-driving cars and online-shopping recommendation systems to name a few.

These examples have one thing in common. They are applications of technology that serves an immediate, or at least foreseeable purpose. For AI in scenarios of this kind, the term *applied AI* (or *weak AI*) has been coined. *Strong AI*, in contrast, is about researching the nature of intelligence itself. An actual (hypothetical) implementation of a strong AI would mean, one would have to build a machine, that is capable of acting like a human being—not just for a limited problem set, but in all of them.

Another term, that is being used more recently, is *AIG*, for *Artificial General Intelligence*. It it is in contrast to what are called *narrow AI* applications, which chess computers and other expert systems would be an excellent example for. AIG sets out, to develop software that can solve and act appropriately in a wide variety of problem fields, without specialising on any particular problem whatsoever. A complete AIG system is supposed to control itself autonomously and to have it's own thoughts and feelings. This has been the original focus of the AI field, before many lost their enthusiasm about it, when it turned out being not as imminently as expected. Even though the advances in narrow AI contribute to it, AIG is more than just an assembly of these. There is a huge number of different AIG projects being worked on, with most of them being in early stages. [GP07]

Cognitive AI, in particular, can be thought of as architectures that implement findings and theories in the fields cognitive science and psychology, as well as the neurosciences, for the sake of proving, if the theories hold against what they promise. Many cognitive architectures share characteristics with or directly implement artificial neural networks.

Looking upon the field of AI from a philosophical point of view, computers seem to deliver enormous potential for learning about how minds work. At the same time, new questions arise. If we knew how cognition works and if we could build machines that simulate it, would these minds be real? And what would the ethical implications be? According to Russel's and Norvig's standard reference *Artificial Intelligence: A Modern Approach* [Rus09], one can distinguish in between two fundamental assumptions. The assumption that computers are able to act *as if* they were intelligent is called the weak AI hypothesis. Thinking that an intelligently acting machine ,in fact, *is* performing cognition, is called strong AI hypothesis.

Putting it differently, weak AI considers computers to be an instrument to research cognitive processes, whereas strong AI considers simulated cognitive processes as actually being cognition.

Trying to figure out, if machines are able to achieve cognition, we often explain by enumerating things that computers can not to. They can not be kind, friendly or have a sense of humour. Neither can they tell right from wrong, fall in love or learn from experience. But what can they do? They are at least partly involved in almost every significant recent discovery in most sciences. They steer cars safer than we ever did. Where *exactly* to draw the dividing line between intelligence and machinery? [Rus09]

The most famous indicator for whether a machine is intelligent or not, is the Turing test. Decades after its formulation, contests regarding to it are still being held—may the best conversationalist win! Many would argue that even an algorithm that passes the Turing test would still not be intelligent. It would trick people into thinking it was, but it would still not be conscious, aware of itself. Moreover, it would not have emotions. [Rus09]

But, would not what we believe is possible and what is not, become obsolete, once we would be able to build something that cognition can not be denied from?

Eventually, this question, might not be as significant as it appears to be. Dijkstra tried to explain that whether something is able to think or not, is first and foremost a matter of language and the interpretation of the word *think*, when he famously said: “The question of whether machines can think... is about as relevant as the question of whether submarines can swim.”

2.2 Psi Theory

This section gives an overview of the basic ideas of the Psi theory. The descriptions are based on *Principles of Synthetic Intelligence PSI: An Architecture of Motivated Cognition* [Bac09] by Joscha Bach.

The theory in its foundations was first described by German psychologist Dietrich Dörner in his books “Bauplan für eine Seele” [?] and “Die Mechanik des Seelenwagens” [?] from 1998 and 2002. Dörner's holistic approach goes beyond classic problem solving and develops a unified model for cognition that implements motivation and emotions. He is convinced that artificial intelligence does not have to focus on different aspects of cognition that have to be looked at separately but that a unified theory will ultimately lead to a deeper understanding of cognition itself.

Dörner's “Bauplan für eine Seele” [?] (“blueprint for a mind”) belongs in the interdisciplinary field Cognitive Science, as with his theory he tries to research psychological foundations with the methods of Computer Science. The theory is an attempt to explain the mind as a machine. It has a notable focus on the emotional system and addresses symbolic and sub-symbolic reasoning. The mind is looked at as a fuzzy and

self-extending causal network structure.

The theory tries to model the human mind as an agent, that is controlled by a structure of relationships and dependencies, that strives to maintain homeostatic balance. Every form of representation of the agent's cognition—may it be percepts, plans or abstractions of space and objects—are represented as directed, hierachic spreading-activation networks. This structure can be visualised as an artificial neural network, which we will call *node net* in the following. Nodes in these networks have gates that may send activation and slots that may receive it. The basic elements a node net consists of are called *quads*. Quads themselves consist of one central neuron and four outer neurons that are called *sub*, *sur*, *ret* and *por* (see figure 2.1). The outer neurons can be connected with the central neurons of other quads and thereby form networks. The four different kinds of outer neurons represent different kinds of relations. *sub* stands for a “has part” relationship, *sur* is the inverse to *sub* and represents an “is part” relationship, *ret* stands for the ordering relationship of succession and *por* represents predecession.

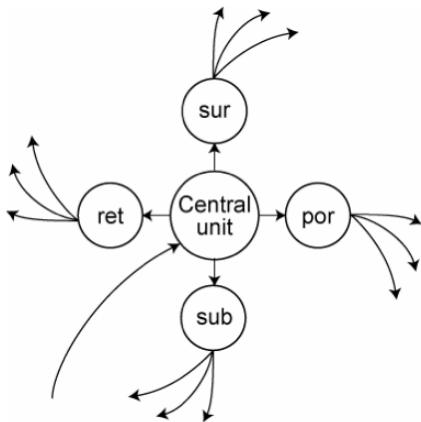


Figure 2.1: Quads like this are the basic representational unit of the Psi theory (taken from [Bac09])

Furthermore each neuron has an activation value, that it receives from and forwards to other neurons. Links in between neurons can be modified to dynamically modulate the forwarded activation value. The modulation of the flow of activation is interpreted as emotions.

There are specific kinds of neurons that are called *sensors* and *actuators*. They represent the in- and output to the outside world. Sensors receive activation if they are triggered by the environment and activated actuators lead to an action in the environment.

The motivation of Psi agents comes from a set of predefined drives. They are divided into physiological (e.g. physical integrity), social (e.g. affiliation) and cognitive drives (e.g. reduction of uncertainty). These drives may signal a specific demand. If the value of a demand changes, pleasure or displeasure signals are sent to the agent. By reinforcement learning agents learn what behaviour leads to the fulfilment of demands.

As node nets may create new nodes and links on their own, the agent forms memory, motives and plans, that help it to select what action to execute next (see figure 2.2). Agents can store representation of their entire percept sequence, but only those that are connected to pleasure or displeasure signals are being kept.

Quads are used to form *partonomies* that represent information about concepts

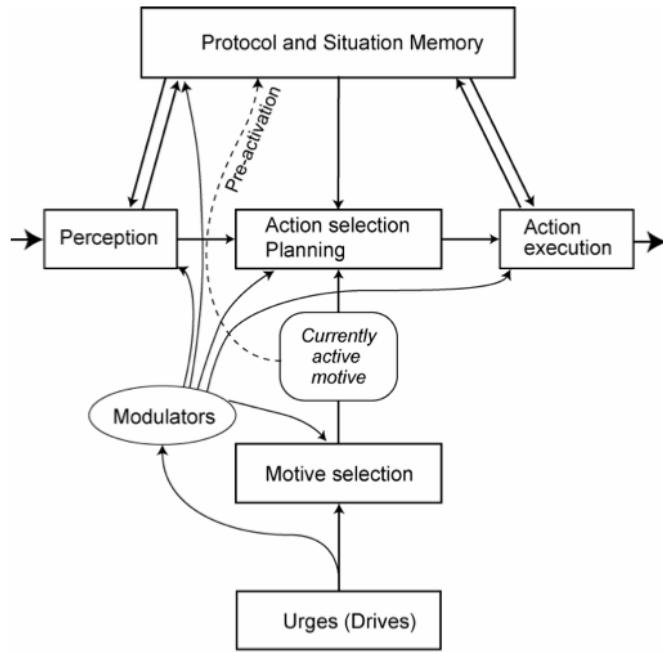


Figure 2.2: The basic architecture of action selection in Dörner's theory (taken from [Bac09])

and relationships. A partonomy is a hierarchical structure that represents a part-whole relationship. For example, the concept “head” would be represented as the assembly of a face and a head. A face furthermore would itself consist of a mouth, a nose, eyes and so forth (see figure 2.3).

Through the combination of partonomies and order relations, all basic types of relationships the agent stores are formed. These could be causal, spatial or temporal relationships in between states of the environment, for example.

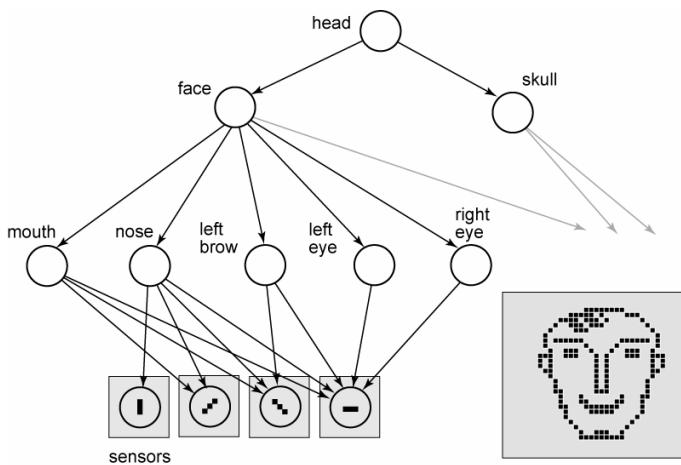


Figure 2.3: The partonomic representation of a Cartoon face (taken from [Bac09])

Psi agents' basic approach to problem solving works by going through a set of subsequent stages. The agent first tries to approach a problem with behavioural routines. If these do not solve the problem, the agent tries to construct a plan, by trying to find a

chain of causal relationships that ultimately lead to a state, that solves the problem. If no plan is found, the agent starts to explore its environment to gather more knowledge.

The goal of implementing the theory to computer software is to find potential inconsistencies and moreover to find out if Psi agents show similar behaviour to that of humans. If they do, the Psi theory could possibly serve as an explanation about how human minds work. For this purpose, a number of experiments—similar to video games—have been concluded, in which both Psi agents and human players had to perform specific tasks. Later on, their behaviours have been compared and searched for similarities and differences.

Joscha Bach adapted the theory to bring it in a contemporary form with slight modifications in his dissertation *Principles of Synthetic Intelligence PSI: An Architecture of Motivated Cognition* [Bac09]. Even though building a conscious machine that thinks and acts as we do is still mere science-fiction, it is this kind of foundational research, that leads us to new ways of thinking about the world, that give us our most significant leaps.

2.3 Agents and Environments

Before we continue to describe concrete implementations and simulations, we define agents and environments according to Russel and Norvig. [Rus09]

An *agent* is defined as anything that can perceive its environment through sensors and conduct actions inside the environment through actuators. The examples given in [Rus09] include a software agent that may receive network packets as sensory inputs and produces output by writing files and sending network packets fits in nicely within this thesis. However, agents can just as well be robots that have cameras as sensors and motors as actuators. Even humans can be thought of as agents that perceive their environment through their senses and act upon it through their muscles.

Russel and Norvig furthermore define the *percept*, as the agent's input and the *percept sequence* as the complete history of what the agent has perceived. Each decision an agent makes is based on the percept sequence.

Additionally, they describe the *agent function*, as the function that maps each percept sequence to an action. The *agent program* is its concrete implementation.

A *rational agent* is an agent that acts appropriately to any given percept sequence. To answer the question about what is appropriate and what is not, they introduce *performance measures* that evaluate the agent's actions. These measures have to be defined by the designer of an experiment. Summarising, the rational agent tries to maximise its performance measure at any given point of time.

Environments have a number of different properties. They can either be *fully observable* or *partially observable*. An environment is fully observable when the agent has access to the complete state of the environment—or at least to everything that is relevant, regarding to the performance measure. As opposed to this, it is partially observable, if the previously mentioned requirement is not given.

An environment can moreover be *single agent* or *multiagent*. An environment is multiagent, when there is at least one other agent, whose behaviour is about maximising a performance measure that depends on the first agent's behaviour. If the agents try to maximise the performance measure of all agents, the environment is called *cooperative*. If agents try to maximise their own and minimise the performance measure of the other agents, it is called *competitive*.

An environment may furthermore be *discrete* or *continuous*. If an environment can be thought of as advancing from one discrete state to the other, it is discrete (and even more so if there is a finite number of states). If the transitions between states are thought of as being continuous, so is the environment.

A *task environment* is the combination of the performance measure, the environment and the agent's actuators and sensors.

2.4 Psi Implementations

Psi has been implemented by different groups at different times. The first implementations are by Dörner and his associates themselves (see figure 2.4). They used Delphi Pascal and developed it for Windows environments. This implementation can still be downloaded ([Psi13]) and runs on Windows 7 installations, for example.

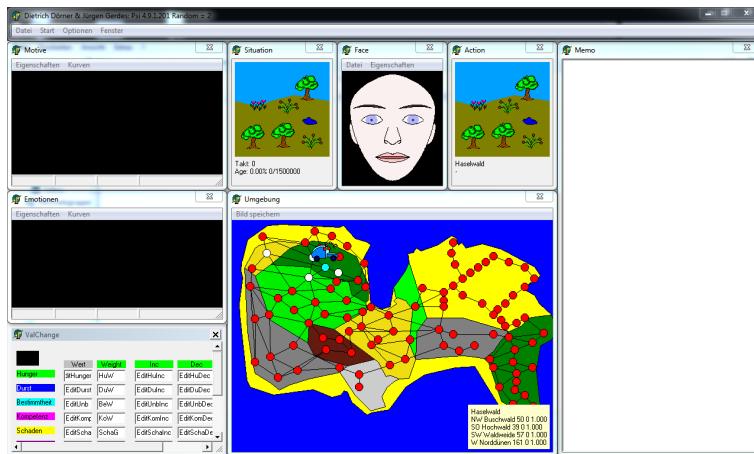


Figure 2.4: Screenshot of Dörner's team's Delphi Pascal Psi implementation

Subsequently, there has been a simple 3D implementation of Dörner's island simulation called Psi3D.

The work on Dörner's team's implementations has not been continued, so Joscha Bach and his associates developed new implementations of Psi. From 2003 to 2009 they built an implementation in Java as a set of plugins for the Eclipse IDE called MicroPsi. It included a graphical editor and a broad experimentation framework. MicroPsi provided a complex and sophisticated simulation component which contained simulated objects in three dimensional worlds—even though most experiments took place on a plane. As a pragmatic approach, different ground types of the simulation world have been stored in bitmap files similar to height-maps. The framework offered complex administrator interfaces and appealing DirectX rendered 3D-views of the scenery as well as a more general world view which interfaced the world by providing clickable objects. Predefined simulation environments included a classic Island and a Mars world (see figure 2.5). Objects in these worlds may be assembled recursively out of other objects and bring their own functionalities. As it has been the case in Dörner's implementation, a viewer for facial expressions has been included—this time as well as a three-dimensional simulation. [Bac09]

Simulated environments proved especially to serve well for the research of collaborative behaviour of multiple agents, for mapping and exploration, image processing as well as for memory and planning. Additionally, some scenarios, that are almost im-



Figure 2.5: Screenshots of the 3D visualisation component in MicroPsi (taken from [Bac09])

possible to implement in the physical world (such as evolving agent populations) could easily be simulated. On the other hand, downsides of simulated worlds include being limited by computing power and the programmer's specifications. [Bac09]

2.5 MicroPsi 2

To ensure broad understandability and to maintain platform independence, MicroPsi has been built ground up again in 2011 and 2012 using more lightweight Python code. What is remarkable about the new implementation called MicroPsi 2 (in the following MicroPsi), is that the simulation is deployed as a web application and the graphical interface is completely rendered inside a web browser, using state-of-the-art internet- and web application technologies. [Bac12]

Even though there have been more complex simulation environments (e.g. 3D-worlds) for previous implementations of Psi architectures, the relatively new version of MicroPsi so far has only two fairly simple ones: a 2D-Island and a map of the public transportation system of Berlin (see figure 2.6). Instead of building a new 3D-world, with this project we set out for something more experimental. More on this follows in chapters three and four.

2.5.1 Module Overview

MicroPsi is written in Python with a minimum of dependencies. Therefore, its modular structure is comparably easy to understand. It is illustrated in figure 2.7. First, one can differentiate in between the `Server` module (or the web interface) and the actual simulation `Runtime` module (also called *core*). In a simple simulation experiment setup, MicroPsi runs three threads: one for the `Server` and, invoked by the `Runtime`, one *world runner* that runs a simulation world as well as one *node net runner* that runs a node net. As these names suggest, the `Runtime` manages both the simulation environments as well as the inhabitant agents (or node net embodiments). They may by design run asynchronously. In fact, the `Runtime` works entirely independently of the `Server` and therefore may just as well be deployed for command line interaction or other GUIs. Furthermore, the `Server` contains a `UserManager` and the `Runtime` a `ConfigurationManager`. [Bac12]

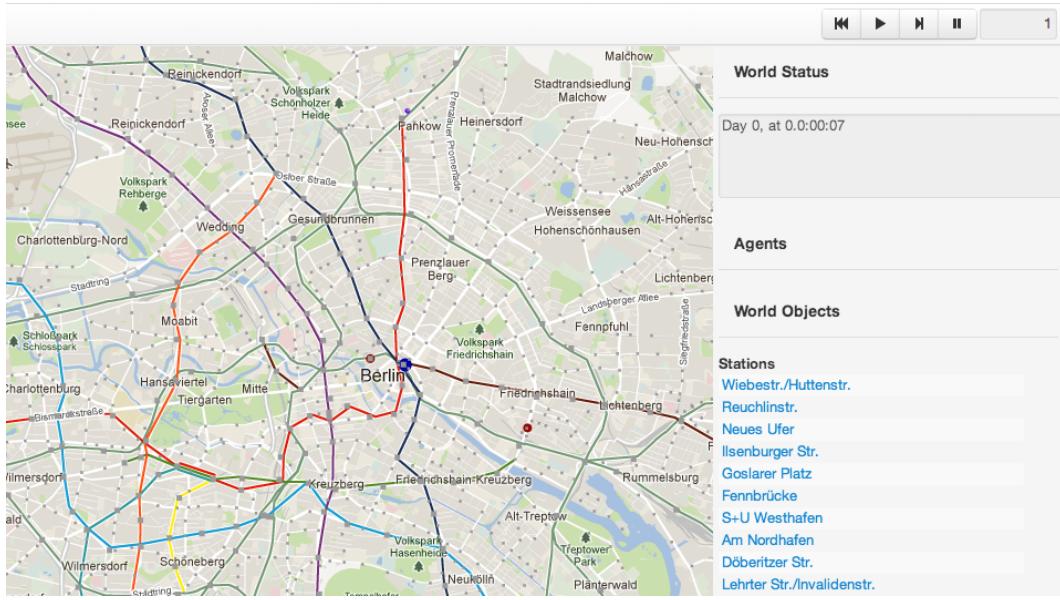


Figure 2.6: MicroPsi’s simulation environment Berlin

The following description is heavily based on [Bac12], where the theoretical foundations can be found in detail.

2.5.2 Module Server

The **Server** renders the GUI and deploys the agent simulation as a web application. It acts as a web server for remote or local access. A client for this application may be any computer with a reasonable up-to-date web browser. Therefore, simulations can be launched from anywhere without requiring any installation. It rests upon the lightweight Python web framework *Bottle* [?].

The web interface is based on HTML as well as Javascript. The communication in between the browser and the simulation is managed via JSON remote procedure calls. Many GUI components of Twitter’s *Bootstrap* library are in use. The graphic renderings (see figure 2.8) utilise the JavaScript graphics library *PaperJS*.

The **Server** communicates with its users through the server API. User sessions and access rights are managed by the **UserManager**.

2.5.3 Module Runtime

In this setup, the **Server** starts the **Runtime**—even though it may also work independently of the **Server**. The **Runtime** component communicates with the **Server** through the MicroPsi API. It manages node nets and worlds.

2.5.3.1 Node Nets

A MicroPsi node net is defined as a set of states, a starting state, a network function, that defines how to advance to the next state, and a set of node types. Data sources and data targets serve as the input for nodes and output towards a world, where a data source is filled with data from the world and data targets are linked to agent actions in the world.

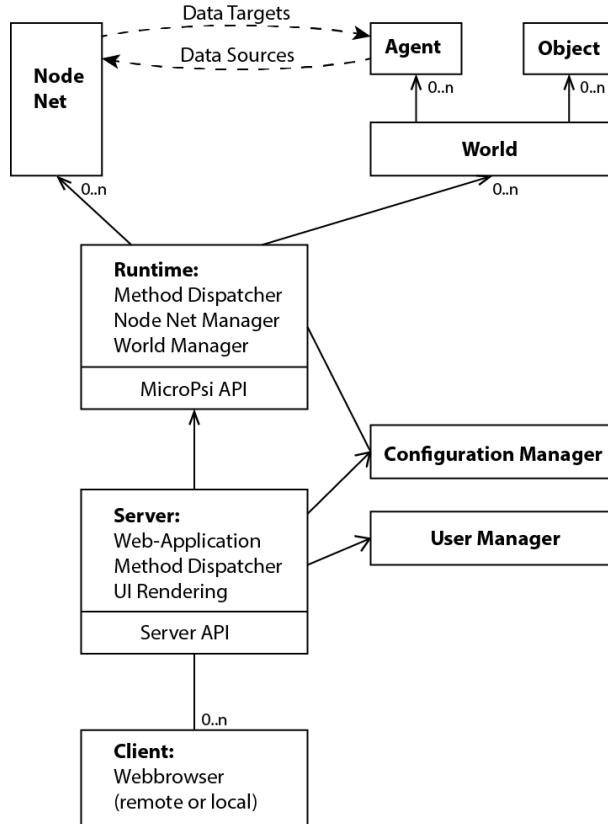


Figure 2.7: The modular architecture of the MicroPsi framework makes it easy to extend. (taken from [Bac12])

According to the Psi theory, nodes may have different types and parameters. They contain gates and slots that send and receive an activation. In most cases, the activation is forwarded from a slot to a gate without further modulation.

Nevertheless, nodes may include functions that enable the creation of new nodes and links as well as procedures for learning and planning. They may be implemented as Python code.

According to the concepts of the Psi theory, MicroPsi defines agents as node nets or to be more specific, hierarchical spreading activation networks. They are an “abstraction of the information processing provided by brain” [Bac12]. Agents can be placed and researched in simulation environments or physically embodied as robots.

As node nets share the relevant characteristics with neural networks, they may enable neural learning paradigms. To store information they can form semantic networks. Furthermore, nodes may contain state machines and other operations, which make it possible to build modularised architectures.

2.5.3.2 Worlds

The simulations worlds are the environments in which we can study our agent’s behaviour. Worlds need to provide a world adapter which functions as the interface in between a node net and the environment. Within the world adapter, data sources and data targets have to be defined carefully, to get a functional and meaningful experiment going. They represent the agent’s sensory input and the motoric output. Sophisticated interconnection of those enables interaction with the environment.

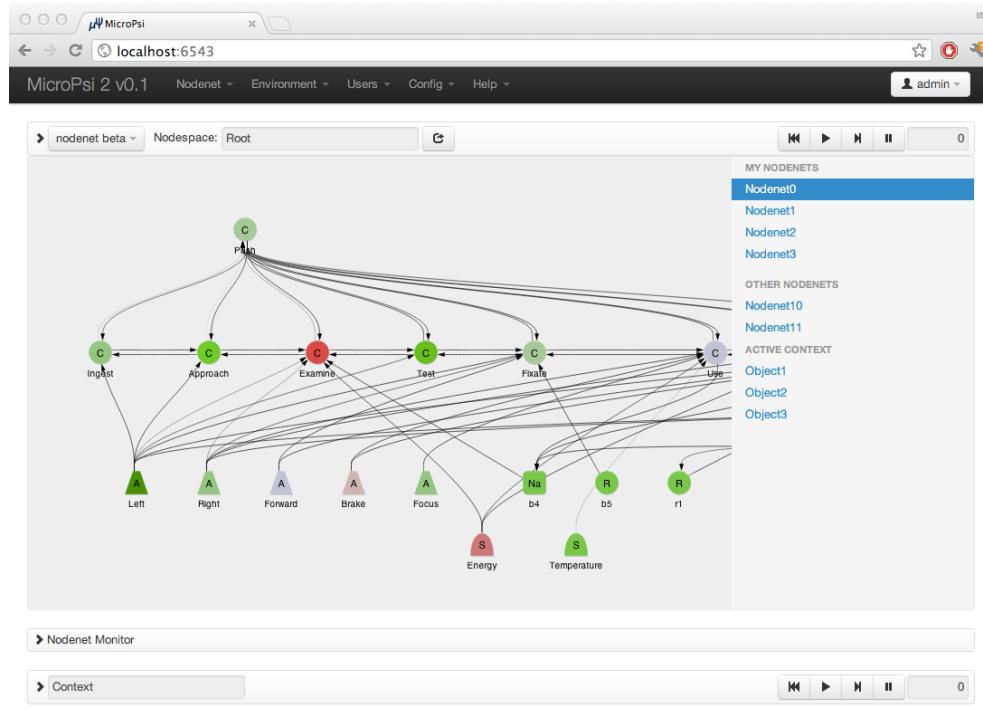


Figure 2.8: The graphical editor is the primary interface to node nets. (taken from [Bac12])

The kind of data the world adapter interfaces, is not specified any further, which gives developers the opportunity to experiment with classic simulation worlds, as well as exotic applications (eg. stock data). At the time of the original development of the MicroPsi, the prioritised application was building a framework for knowledge representation.

Objects

Worlds contain objects. Objects may be anything that could be interesting for a simulation, and that needs some kind of integrated logic. Light sources or collectable resources are a common example. Objects may contain a set of functions and states, but at least one function that determines how the object advances and reacts to changes while moving through a simulation cycle. A comprehensive world function calls every object function for each simulation step of the world.

Agents

Agents are objects that are connected to a world adapter which makes them controllable by a node net. The object that incarnates the agent is best thought of as the agent's body. The function that advances the agent checks for input from the node net and outputs to it.

2.6 Summary

This chapter provided the necessary foundations regarding AI for this Thesis. It introduced Artificial General Intelligence as a subSet of AI research and presented its main ideas and goals, as well as regarding foundational philosophical questions. Then, the

Psi theory by Dietrich Dörner has been described briefly. An overview of implementations of the Psi theory has been given. The latest framework dedicated to the Psi theory, MicroPsi 2, has been described and it's architecture outlined. Knowing about MicroPsi 2's modules is a necessary foundation for the additions to it that are described in chapter 4.

Chapter 3

Minecraft

To understand how and why the video game Minecraft [?] is expected to be an interesting simulation environment for artificial general intelligence, necessary background information, about what makes Minecraft different from many other games, has to be given first. To furthermore work with it, the essential components of the game's architecture have to be described. This chapter provides a roundup of Minecraft's history and a brief explanation of the basic game mechanics. They are followed by an overview of Minecraft usages that exceed the original game's purpose as well as an explanation, of why Minecraft is suitable as a simulation environment. Eventually, related software projects, that have been used for this thesis are introduced and described.

3.1 Overview

The story of Minecraft has many interesting sides, but first and foremost it is a story of immense, unexpected success. To give an overview of it, this section delivers insights into its central ideas, as well as a general look upon the developer's practices. Moreover, the game's reception is considered, as the game has become unusually popular. The descriptions are based on Sean C. Duncan's article *Minecraft, beyond construction and survival* [Dun11].

3.1.1 Main Ideas

When Markus ("Notch") Persson built and released the first public version of Minecraft, it soon became apparent that his creation resonated with many people. The simple concept of a world entirely consisting out of standard sized building blocks, which the player can create, destroy and relocate one-by-one, enabled many gamers to employ their creativity, explore the Minecraft world and test out its possibilities and boundaries.

Looking at the game for the first time, it is impossible to oversee the primitive appearance of its graphics engine. Everything in the game is made out of blocks. May it be the leaves of a tree, dirt, people or the clouds. Even the sun does not appear round, but as a cubic object. One can assume that the obvious connotation to LegoTM is more than a coincidence, as it immediately gives the player a clue, about what kind of possibilities for building and creating the game might deliver.

Another interesting aspect about Minecraft is the procedural semantics the game world is generated with. Trees in Minecraft, for example, may share a similar structure that consists of a trunk and leave-covered branches spreading out fractally, but the

particular characteristics of each tree are generated randomly. This makes a Minecraft world somewhat more realistic than those of many other videogames.

Minecraft strives to be the opposite of a game with a strict game-flow. Instead, it enables the player's creativity and joy of creation.

3.1.2 Development Process

The development process of Minecraft is very different from those of most other well-known games. Instead of building the game over months or years to finally release a (more or less) finished version to the public, Minecraft is being designed iteratively. In fact, Persson released the first version only a week after he started working on the game. Ever since, new features have been added, and new versions of the game have been released frequently. Another defining aspect of the game's development studio Mojang's approach towards the development of Minecraft is the involvement of the game's players. Since the earliest stages, plans for the game's future have been discussed openly, and many times the players' voices have been heard. Persson himself locates his approach in the field of agile software development. [Dun11]

3.1.3 Reception

The game has attracted great attention ever since its first release in 2009. Since copies of the game could be obtained commercially for the first time, its different versions sold more than 26 million times—with the PC version priced at 19.95€, for example. It should be noted that Minecraft's development studio Mojang is a so called *indie game developer*, that is not associated with any classical game publisher, but distributes copies of their game exclusively via their own website.

Minecraft can certainly be called a world wide phenomenon. The game's developers received a number of prizes for their creation—not to mention the millions of dollars that have been earned in Minecraft sales. With its huge success, Minecraft has proven clearly that it compelled to many players.

Unlike any other game, Minecraft fosters the creativity of its players. Bringing the construction component to one extreme after the other, well publicised player creations include a life-size model of the USS Enterprise-D as well as a fully functional arithmetic unit. Videos of roller coaster rides through complex constructions increasingly show characteristics of an art form themselves—regarding to the amount of views on youtube, not less than other, well established formats.

3.2 Game Mechanics

Now knowing about the game's background, in this section we will continue with descriptions of the basic concepts of the game, that regard to our AI interface.

Before a new Minecraft game begins, it procedurally generates a three-dimensional world. The player is then dropped to a spawn point and is now facing a scenario and a minimal game interface, lacking any kind of instruction or advise about what to do next. Soon, the procedurally generated world drives most players to explore the highest mountains and deepest caverns. As the player collects resources and crafts items, resources harder to reach and items more complex to craft become available. [Dun11]

Although the game can be downloaded and played as a single packet of software, many scenarios of playing the game consist of running a Minecraft server software, as well as one copy of the client software for each player. It is possible to mimic the official

client by implementing the reverse-engineered client-server-protocol and therefore build artificial players that way.

The structure of Minecraft worlds and the client-server-protocol are described in the following. The descriptions are based on the regarding articles of the *Minecraft Wiki* [mcw13a].

3.2.1 Minecraft Worlds



Figure 3.1: A grass block
(CC-BY-3.0 Mojang AB)
[ima13]

The most significant concept in Minecraft is the block (see figure 3.1). A block is a cube which sides are, compared to the player, roughly one meter long. Every Minecraft world is built up entirely out of them. Therefore, worlds can be thought of as three dimensional spaces filled with voxels. [BB] There are many different types of blocks (or materials, they consist of) and they share their single size, which converts to the basic distance unit of Minecraft.

A chunk (see figure 3.2) is a segment of the Minecraft world that is 16 blocks long, 16 blocks wide and 256 blocks high (or deep) and therefore consists of up to 65,536 blocks. [mcw13b]

"The player" (see figure 3.3) is what the playable game-character in Minecraft is called. It is usually displayed as a human. Also, a Minecraft world has a day-night-cycle with 24 Minecraft hours converting to 14 minutes by default.

The game itself has no predefined goals. Players can walk around, discover the created world (see figure 3.4 1) and collect resources by *destroying* blocks, with the generated resources equaling the block type. They can combine several resources to *craft* items. For example, a player can destroy the blocks that represent a tree (see figure 3.4 2). The gained *wood* resource could then be used to craft a wooden pickaxe (see figure 3.4 3), which could then be used to dig into the ground more efficiently (see figure 3.4 4) to then *mine* more rare resources, like iron or gold.

There exist different game modes. The original *survival* mode adds monsters that attack the player at night. What solutions to survive the player comes up with (e.g. building shelter or fighting the monsters) is left up to him or her.

In *creative* mode, the player is not being attacked by monsters, has the ability to fly and is given instant access to unlimited resources.

The mode of a Minecraft world does not effect the functionality of this project.

3.2.2 The Client Server Protocol

Minecraft's client server protocol is not publicly documented by the developers themselves. However, the modding community gathered comprehensive knowledge and understanding of its structure (probably by using reverse engineering techniques). The protocol is based on packets.

The NBT (Named Binary Tag) is a data format that has been introduced by Persson to transfer chunks of binary data together with small

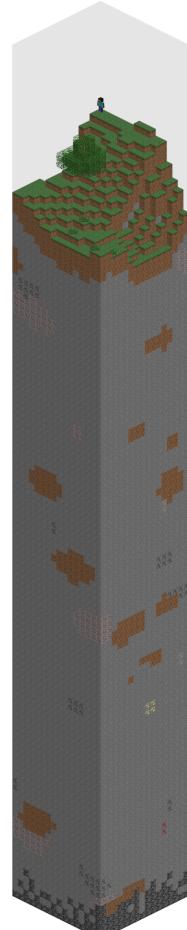


Figure 3.2:
A chunk
(CC-BY-3.0
Mojang
AB)
[ima13]



Figure 3.4: Screenshots of Minecraft's basic game mechanics

amounts of additional data. Its main use is to store Minecraft worlds.

Packets are either *server to client*, *client to server* or *two-way* and begin with a *packet ID* byte. The structure of the packet's payload depends on its packet ID.

To give an example of one of the easier packets, the *Client Position* packet is fairly straightforward (see table 3.1). It is exclusively send from clients to servers and starts with its packet ID (as every packet does), followed by the X- and Y-coordinates as doubles, the stance value as a double (which is used to modify the player's bounding box), another double for the Z-coordinate and eventually a boolean that describes if the player is on the ground or not. [pro13]

Field Name	Field Type	Notes
Packet ID	Byte	0x0B
X	double	Absolute position
Y	double	Absolute position
Stance	double	Used to modify the players bounding box
Z	double	Absolute position
On Ground	boolean	Derived from packet 0x0A

Table 3.1: The structure of the packet "Player Position (0x0B)" [pro13]



Knowledge of this data structure is already sufficient to move around in the Minecraft world. To go forward, one has to figure out the players current position, calculate the absolute coordinates of the destination of the movement in regard to it and send a Client Position packet with these coordinates via a TCP socket to the server. If the destination is not more than 100 blocks away from the origin of

Figure 3.3: The player
(CC-BY-3.0 Mojang AB)
[ima13]

the movement, the server accepts the movement. In the official Minecraft client, a player's movement from one point to the other is rendered with a walking animation.

Other than movement, packet structures are defined for every aspect of the game. May it be the initial handshake, the creation or destruction of blocks or activities of other player- or non-player-characters.

This protocol is what a custom Minecraft client needs to speak.

3.3 Minecraft as a Gaming Platform

With more and more applications that move beyond the original gameplay of construction and survival, Minecraft seems to shift gradually towards being not only a game, but also a gaming platform, that many other people use to implement their own ideas. Inside Minecraft worlds, players have created games for educational purposes, for example. A group of students of the Miami University developed a game called Circuit Magic that is committed to teaching the players about logical circuits. The famous Minecraft Teacher, who implements Minecraft in primary school classes, is another good example. Following in his footsteps, "Massively Minecraft" is a community of teachers who share their thoughts and experiences with doing the same. Just as well, games have been developed for the purpose of artistic expression. Chain World, the game that tried to simulate religion, is an exciting example. Not only do players implement their engagement inside Minecraft. The internet is full of Minecraft content, be it videos on Youtube, online tutorials or other community based knowledge resources, headquartered at the Minecraft Wiki¹. [Dun11]

As user-generated content has been an essential part in the development of Minecraft, in a process that reminds one of a self-fulfilling prophecy, more and more people explore its possibilities to experiment with new applications. As of today, Persson and Mojang seem to favour and approve new uses of their creation. [Dun11]

3.4 Suitability of Minecraft as a Simulation Environment

There are a number of reasons why using Minecraft as a simulation environment could be useful and lead to interesting results.

First, the game itself is easily accessible. It is developed using Java (for both the client and the server software) and, therefore, up to a certain extend, platform independent. The desktop computer version is being sold for Windows, Mac OS X and Linux devices. There are official ports for Android, iOS, Xbox 360, the Raspberry Pi and a version for the upcoming gaming console Xbox One is announced. The desktop versions are priced at 19.95 €, which make it affordable to a large audience.

The game itself already has an enormous fanbase. It is (like most videogames) especially popular among teenagers. Minecraft being loved by so many people could benefit this project, in terms of leading to increased attention.

The game's developer has proven many times that it acts generously towards other developers, when it comes to the creation of game modifications and content that builds upon or changes original Minecraft intellectual property. As mentioned before, Mojang is not restrictive towards users doing all kinds of things with their creations. This led to the availability of a fairly complete community-sourced documentation and

¹<http://www.minecraftwiki.net/>

explanation of virtually every aspect of the game—including its software architecture, data structures and protocols. This is useful for this project, as chances are low that they will have anything against using Minecraft for AI simulations in the foreseeable future. To the contrary, the game’s AI creator Jon Kagström provided advise for this project

The Minecraft world, with its logic, semantics and functionalities, offers possibilities for an AI to prove being able to interact with the environment—in primitive ways (e.g. moving around), as well as with increasingly complex tasks like building, collecting resources, crafting items and interacting appropriately with both well-disposed and hostile other entities.

The semantics of the gameworld share characteristics with the real world. Moving through a Minecraft environment, one quickly realises that the game has generated different biomes (eg. forest or tundra). Also, trees, rivers, mountains and ore veins are neither hard-coded nor appear completely random, but are generated procedurally and their structure appears to be (somewhat) fractal.

Using Minecraft as a simulation environment will give Psi agents possibilities to show off, what kind of sophisticated behaviour they are capable of.

3.5 Related Projects

Minecraft’s success inspired many other projects, including artificial players and a number of Minecraft-like games in a wide variety of programming languages and environments.

3.5.1 Minecraft Bots

There exist many projects that could be considered Minecraft *bots*. One has to differentiate in between two types. On the one hand, there are those that mimic an entire client software and facilitate communication with the server on the default client software’s behalf. On the other hand, there are bots which are modifications of the original client (or server) software and usually add non player characters—like animals and other non-human creatures—to the game. The code is usually injected through one of the popular *modloaders* (e.g. Minecraft Forge).

One example (and probably the most advanced one) for an entire bot framework that replaces the client, is *Mineflayer* [git13]. It has a high-level abstraction of the environment (eg. entity knowledge and tracking) and is written in JavaScript using node.js. However, it has not been used for this project, because a Python implementation was aimed for, as it is explained in the next section.

Opposed to Mineflayer, an example for game modification bots are the *Cube-bots* [mcf13]—fan-made non-player characters that aim to help Minecraft players with mundane tasks.

3.5.1.1 Spock by Nick Gamberini

Developed by Nick Gamberini, Spock² is an open-source bot framework (and as such also a Minecraft client) written in Python. It has been chosen to become an essential part of this project for two reasons: being written in Python it painlessly integrates in

²<https://github.com/nickelpro/spock>

the existing MicroPsi code and the absence of dependencies (with one exception) leave the code understandable and easy to deploy.

3.5.1.2 Protocol Implementation in Spock

The Minecraft protocol implementation in Spock is straightforward. The necessary data structures are stored separately and can be accessed globally (see figure 3.5).

```

1 names = {
2     0x00: "Keep Alive",
3     0x01: "Login Request",
4     0x02: "Handshake",
5     0x03: "Chat Message",
6     ...
7
8 structs = {
9     #Keep-alive
10    0x00: ("int", "value"),
11    #Login request
12    0x01: (
13        ("int", "entity_id"),
14        ("string", "level_type"),
15        ("byte", "game_mode"),
16        ("byte", "dimension"),
17        ("byte", "difficulty"),
18        ("byte", "not_used"),
19        ("ubyte", "max_players")),
20    ...

```

Figure 3.5: An excerpt of Spock's data structures for packet-IDs and -structures

These structures are used to parse each packet appropriately (see figure 3.6).

```

1 def decode(self, bbuff):
2     #Ident
3     self.ident = datautils.unpack(bbuff, te')
4
5     #Payload
6     for dtype, name in ta.structs[self.ident][self.direction]:
7         self.data[name] = utils.unpack(bbuff, dtype)

```

Figure 3.6: Spock's function for decoding packets

3.5.2 Minecraft Clones

Other interesting projects include Skycraft [sky13], a Minecraft-like browser game based on WebGL.

A particular project that has the same name as the original game that inspired it is **Minecraft** by Michael Fogleman (see figure 3.7). It is a simple Minecraft clone in under 600 lines of Python and has gained some popularity on reddit [fog13a] and Hacker News [fog13b].

It is comparably easy to understand and modify and has been used for the visualisation component of this project. It is based on the Python multimedia library

Pyglet [pyg13].

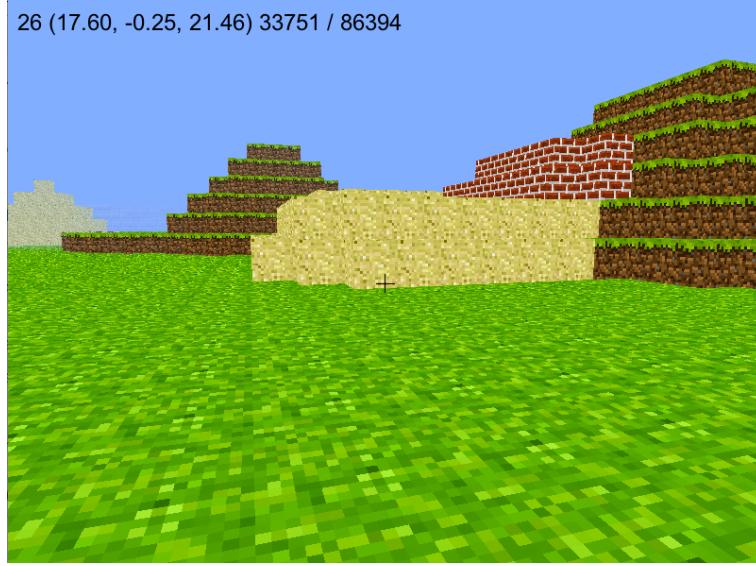


Figure 3.7: A screenshot of *Minecraft* by Michael Fogleman

3.6 Summary

An overview of Minecraft, its history and its basic concepts, as well as other software in its periphery, have been presented in this chapter. Minecraft is a complex, yet easily accessible virtual world. It is in constant development and new features are added regularly. It has a massive fanbase and a huge community around all kinds of game modifications. The most important concepts of Minecraft for this thesis are blocks, chunks and also the client server protocol. They will all be used in the next chapter. As a basis for the implementation of the Minecraft interface for MicroPsi, the bot framework Spock by Nick Gamberini has been chosen to extend MicroPsi. Furthermore, *Minecraft* by Michael Fogleman has been chosen to serve as a foundation for the visualisation component. Their implementations will be presented in detail in the next chapter.

Chapter 4

Minecraft as a MicroPsi 2 World

The objective of this project is to build and test an interface in between MicroPsi and Minecraft, so that a Minecraft world can be used as a simulation environment for experiments within the MicroPsi framework, which will act as an artificial player. Since there exist open source projects that perform many of the tasks required for this goal, **Spock**, the Python Minecraft bot framework by Nick Gamberini, and **Minecraft**, the Python Minecraft clone by Michael Fogleman, have been chosen as a foundation. To make the new simulation environment monitorable, the latter has been used to implement a visualisation of the Minecraft world in the web interface. The following section gives an overview of the implemented modules.

4.1 Overview

The modular architecture of MicroPsi allows it to add new simulation environments (or worlds, as they are called in MicroPsi) fairly easily. To communicate with a MicroPsi node net, a world needs an interface, which is called *world adapter*. The *world adapter* has to define *data sources* and *targets*. It fills the sources with data from the world and writes the targets to the world. The node net does the opposite: it reads from the sources and writes into the targets. This enables a feedback loop in between the world and the node net. Furthermore, the *world adapter* provides a step function, that advances the world and is called by the MicroPsi world runner frequently.

Looking at the Minecraft side, communication with a Minecraft Server typically requires a constant flow of data packets going in and out. Most third party clients, including bots, facilitate their own event loops. To add a Minecraft world to MicroPsi, the demands of both sides have to be met.

The contributions of this project are divided into the three modules `minecraftWorld`, `minecraftClient` and `minecraftVisualisation`. The resulting architecture is displayed in figure 4.1. The `minecraftClient` manages the communication with the Minecraft server, provides convenient functions and data structures for sending and responding to packets and stores and regularly updates a simple representation of the environment data it receives from the server. The `minecraftVisualisation` module generates 3D-Images that display the current state of the Minecraft environment, based on the data it receives from the `minecraftClient`. What ties it all together is the `minecraftWorld`. It provides a step function that advances both the `minecraftClient` and the `minecraftVisualisation` and is called itself by the world runner of the MicroPsi framework. Furthermore, it defines and updates the *data sources* and *targets*.

The `minecraftVisualisation` module can be exchanged or cut off completely very

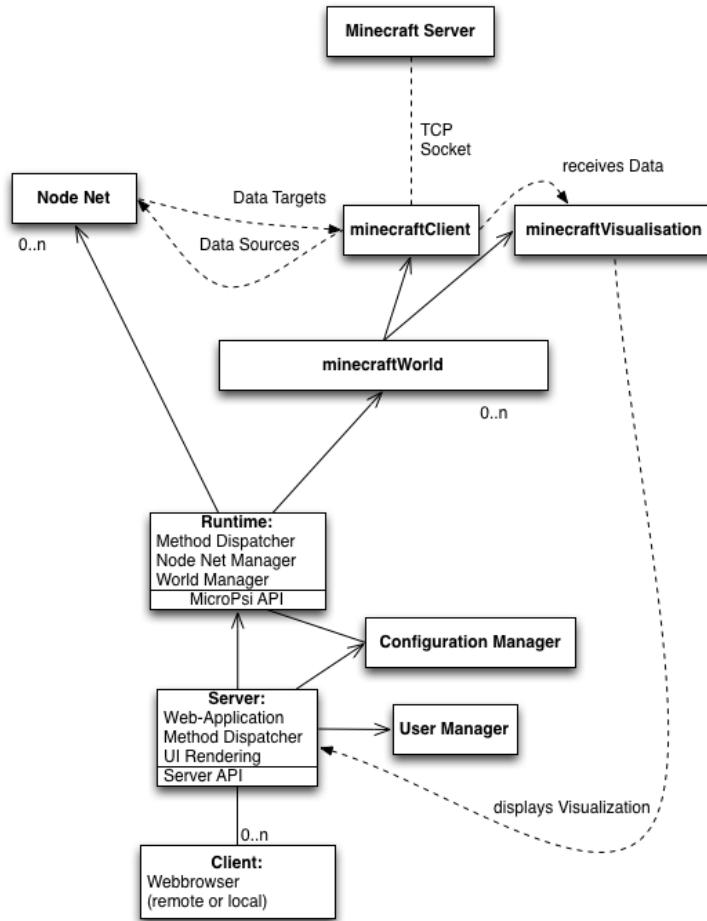


Figure 4.1: The new architecture of MicroPsi with the Minecraft interface

easily, as no other modules depend on it. Instead of the visualisation, a placeholder image can be displayed in the web interface, which does not effect the functionality of the simulation. The **minecraftVisualisation** module itself depends on the data structures of the **minecraftClient**, though. This means, exchanging the **minecraftClient** would require adjustments of the **minecraftVisualisation**, to still function as intended. The same holds true for the *data sources* and *targets* in the **minecraftWorld**.

4.2 Using Spock as the **minecraftClient**

As mentioned above, the purpose of the **minecraftClient** is to manage the communication with the Minecraft server and to provide a representation of the agent's environment.

The calculation of the simulation environment, does not take place in MicroPsi itself, but on a regular Minecraft Server. Instead, **Spock** is integrated into MicroPsi and represents the simulation world towards it. **Spock** (in the following the **minecraftClient**) communicates with the Minecraft server via the client server protocol and provides data that can be used as *data sources* for the world adapter and translates the data from the *data targets* to actions in the simulation environment. That way, to MicroPsi it looks like the **minecraftClient** is the simulation environment itself, where, in fact,

it is the interface to the game world server.

The original event loop of the bot framework had to be dissolved and rebuilt as an `advanceClient` function that is called as a part of the world adapter's step function. The event-loop and -handling of the `minecraftClient` had to be slightly modified to work with MicroPsi. It should be noted that the frequency, with which the framework steps the bot, has to be at least chosen high enough (about 10/s), so that Spock is able to regularly send the necessary keep-alive-signals, to not get kicked from the server.

For every iteration of the event loop, the `minecraftClient` reads incoming data from the socket, dispatches the read packages appropriately and eventually checks the MicroPsi *data targets* to perform an action—if necessary (see figure 4.2). Note that the described event loop of the original Spock is not a loop anymore, but each iteration is invoked as a part of the world adapters step function.

Eventually, useful data sources had to be picked and a system of data targets and their translation to actions had to be implemented. In most cases, performing actions means to let Spock send a particular set of packets to the Minecraft server.

4.2.1 Overview of the `minecraftClient`

The `minecraftClient` is heavily based on Spock, that has been developed as an educational project. The scope of this project was to build a pure Python Minecraft bot framework without dependencies. This has been achieved with one minor exception: if one would like to connect the bot to an official Minecraft online server, the packets have to be encrypted using the Python cryptography library PyCrypto.

The `minecraftClient` runs as a part of the world runner thread. It consists of several classes.

The main class, `minecraftClient`, holds references to instances of the classes `BoundBuffer`, `World` and `Packet`. Furthermore, basic data structures are stored in the files `cflags` and `mcdat`.

The class `BoundBuffer` is an implementation of a buffer that matches the particular needs of sending and receiving Minecraft packets.

The class `World` holds the internal representation of the gameworld. It brings functions, data structures and subclasses to represent chunks and to obtain information about which block sits where. In its heart, it manages a dictionary, that stores chunks as binary data.

The class `Packet` represents a Minecraft packet and brings functions to encode and decode it (being read from a `BoundBuffer`) to get to its payload or to be able to send it to the Server.

The file `packet_handlers.py` provides a class for each packet that the client is supposed to deal with by default. The following listing gives an example.

The dictionary `cflags` provides the Socket codes. The file `mcdat.py` provides dictionaries for the datatypes, block type codes, packet names and structures of packets of the Minecraft protocol.

In the file `nbt.py` several classes for dealing with the NBT file format exist.

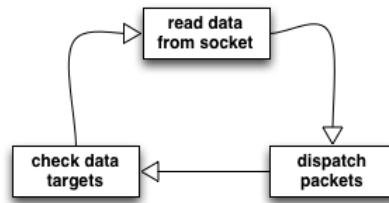


Figure 4.2: The `minecraftClient`'s event loop

```

1 #Chunk Data - Update client World state
2 @phandle(0x33)
3 class handle33(BaseHandle):
4     @classmethod
5     def ToClient(self, client, packet):
6         client.world.unpack_column(packet)

```

Figure 4.3: By default, this function handles "Chunk Data (0x33) packets.

The file `timer.py` contains the classes `EventTimer`, `TickTimer` and `ThreadedTimer`, which provide measures for the timing of the Minecraft world.

The client sets up a socket to the server, starts off with a handshake to from then on facilitate packet based communication with the server.

4.2.2 Extensions to the `minecraftClient`

It was aimed for to extend the original client in a way that it would fit in nicely as a simulation world for MicroPsi.

A reference to the new class `PsiDispatcher` has been added to the `minecraftClient`. Its purpose is to check the `minecraftWorld`'s data targets frequently and invoke appropriate actions, if necessary. The following figure provides a simplified example of the resulting architecture. The `PsiDispatcher` is called as a part of the `advanceClient` function. It checks each available data target, and if necessary invokes an appropriate action (eg. sending a packet). Listing 4.4 is an example for the data targets that indicates that the agent is supposed to move one block towards the direction of the x-axis.

```

1 #check for MicroPsi input
2 if self.client.move_x > 0:
3     self.client.push(Packet(ident = 0x0B, data = {
4         'x': (self.client.position['x'] - 1) / 1,
5         'y': self.client.position['y'] / 1,
6         'z': self.client.position['z'] / 1,
7         'on_ground': False,
8         'stance': self.client.position['y'] + 0.11
9     }))

```

Figure 4.4: Checking for activation of the move-x data target and invoking the according action.

The event loop has been replaced by the function `advanceClient` that is called for every simulation step. For example, one iteration could mean receiving and storing new block data from the Minecraft Server, sending default responses to the received packets, checking the MicroPsi data targets for moving and if necessary send a "Player Position" packet to the Minecraft server.

4.3 Module `minecraftVisualisation`

That being said, the other important part of this project is the visualisation component. It contains classes that provide an interface to the OpenGL context that is the

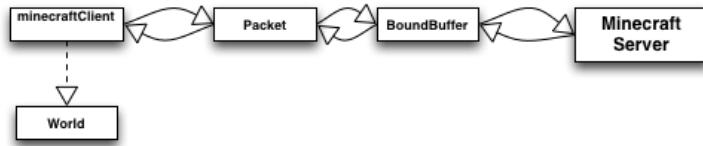


Figure 4.5: An overview of the most import classes of the `minecraftClient` in respect to the communication with a Minecraft Server

visualisation. This section gives an overview about what data sources the visualisation uses and how it could be replaced or extended. Inside the `minecraftWorld`'s step function, the visualisation module is called to generate a 3D model of the Minecraft world and the agent within. There are two main reasons for this. The first reason is that the agent's behaviour within the simulation environment is supposed to be visually monitored from the MicroPsi web interface—in a both efficient and pleasurable manner. The second reason is that the image data is supposed to be processed by the node net as a data source itself in the future.

The visualisation component reads from the `minecraftClient`'s internal gameworld representation to generate the 3D model. This means that, from pure Minecraft world data, a 3D-visualisation has to be generated from within the MicroPsi Python code. It should contain a perspective that gives a good overview over the bots environment and forward it to the web interface. The visualisation is in its core based on “Minecraft” by Michael Fogleman.

Specifically, the representation of the chunk, the agent is located in, is fetched and for each solid cube in this chunk, a corresponding cube is rendered within the visualisation using Pyglet’s OpenGL abstraction (see figure 4.6). Each block gets textures according to its type. The implemented format for the textures is compatible to the widely available Minecraft texture packs. That way, the visualisation’s look can be changed completely within seconds. The resulting images are exported as JPEG files. Then, they are displayed in the web interface. A refresh rate of six or more images per second creates the impression of a video stream.

The module `minecraftVisualisation` consists of two classes. The class `Window` inherits from `pyglet.window.Window` and therefore initialises the OpenGL context. For every call of the `advanceVisualisation` function, it updates the 3D model according to the world representation inside the `minecraftClient` and renders it. A PNG snapshot of the frame buffer is generated and displayed in the web interface.

The class `Model` is also used as a part of the `advanceVisualisation` function. It contains functions for adding and removing blocks from the OpenGL canvas. The textures for the blocks are loaded from PNG images, stored as `pyglet.graphics.TextureGroups` and are assigned to vertices when needed.

Similar to Spock, Minecraft by Michael Fogleman implements its own event-loop and -handling. Again, the event loop had to be disassembled and rebuilt as a part of the world adapter’s step function, which advances the visualisation with every step.

4.4 Module `minecraftWorld`

The world adapter contains two classes. The class `MinecraftWorld` inherits from the MicroPsi class `world` and provides the assets for the web interface, an `init` function and the `step` function. The class `Braitencraft` holds of the actual world adapter.

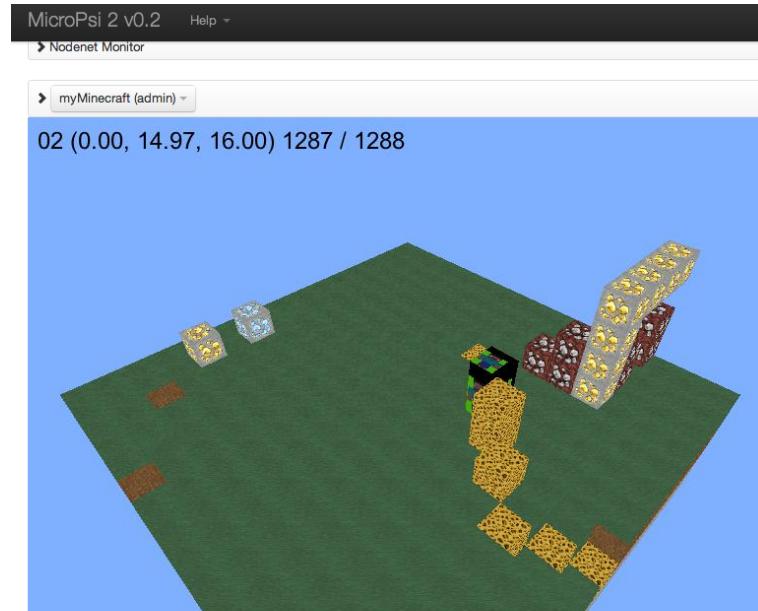


Figure 4.6: An image generated by the `minecraftVisualisation` component

In it, the `data sources` and `targets` are defined as dictionaries and an update function advances the life of the client.

Data Targets and sources

As a proof of concept, four data sources have been defined, that symbolise sensors, that detect, if a block of diamond in the current section is either before, behind, left or right of the client. The data targets are filled, as it is outlined in figure 4.7.

On the other hand, data targets have been defined, that represent actuators for moving forwards, backwards, left and right.

4.5 Case Study

To explain the usage of this implementation, an example is given in the following. An agent is placed in a Minecraft world. It shall contain the four previously mentioned sensors that detect if a block of diamond is around the agent and that point to the positive and negative x- and y-axes. These sensors detect if a diamond ore block has been placed in the current chunk, in the direction they are facing. The agent furthermore consists of four actuators that make the agent move towards each of the same four directions. Next, the sensors get connected to the corresponding actuators, which will lead to an agent being attracted towards a placed diamond block.

To conclude the experiment, a chunk in a Minecraft world is set up that way, that it contains only a plane, without other obstacles, that the agent might move around on freely, as well as a diamond block that is placed at its center.

Then, the agent is placed in one of the chunk's corners. We furthermore define that the activation of the sensors equals the distance to the diamond block in the regarding direction. The data sources get filled through searching in the entire section,

as figure 4.7 displays.

```

1 for y in range(0, 16):
2     current_section = current_column.chunks[int((bot_block[1] + y - 2) / 16)]
3     if current_section != None:
4         for x in range(0, 16):
5             for z in range(0, 16):
6                 current_block = current_section['block_data'].get(x, _block[1] + y - 10 / 2) \% 16), z
7                 if current_block == 56: #Diamond Ore
8                     diamond_coords = (x + x_chunk * 16, y, z + z_chunk * 16)
9                     self.datasources['diamond_offset_x'] = bot_block[0] - coords[0] - 2
10                    self.datasources['diamond_offset_z'] = bot_block[2] - coords[2] - 2
11                    self.datasources['diamond_offset_x_'] = ((bot_block[0] - coords[0]) * -1) - 2
12                    self.datasources['diamond_offset_z_'] = ((bot_block[2] - coords[2]) * -1) - 2

```

Figure 4.7: Searching the current section for a diamond block and, if it is found, filling the concerning data sources with the distance towards it.

In every step of the world, the client will check the data targets and therefore move towards the diamond. Once it gets to a distance towards the diamond, that is below a defined threshold, the sensors will stop sending activation, and the agent will stop moving towards it. Note that if the node net runner and the world runner run asynchronously, there might be a delay in the shift of behaviour of the agent. Hence, it is advised, to run them synchronously or at least with a timing close to each other.

In the node net editor, we set up the corresponding node net, as it can be seen in figure 4.8. Now, the world writes the activation from the sensors to the actuators in every iteration of the step function.

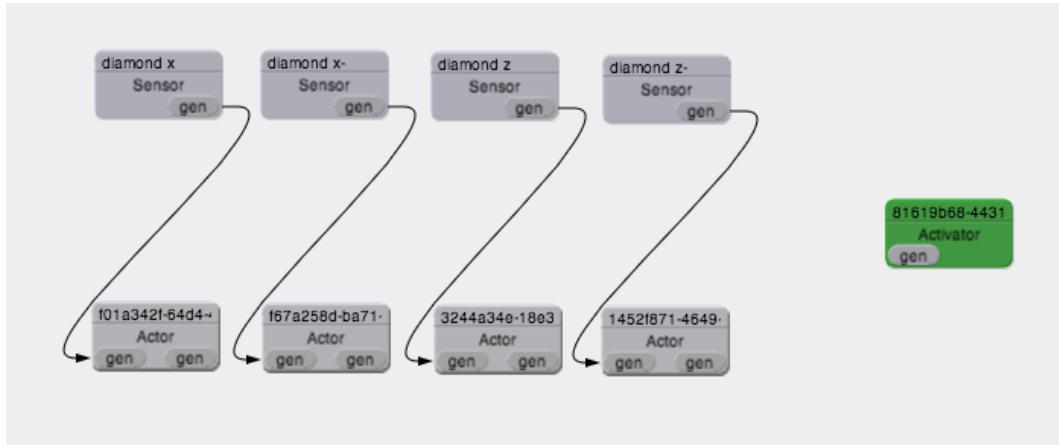


Figure 4.8: The node net set up for this experiment

If we start a simulation like this, the nodes of the sensors that point to the diamond light up green and their activation is forwarded to the actuators. The bot moves towards the diamond until it is closer than the sensor's threshold, to not detect the diamond anymore—for this experiment the threshold has been set to two blocks. Then it stops moving (see figure 4.9).

4.5.1 Evaluation

Now, that a proof-of-concept experiment has been concluded, the implementation can be evaluated. The experiment obviously shows that an interface in between a Minecraft

world and the MicroPsi framework has been implemented. In the experiment, the agent concludes the appropriate action, if a node net actuator node receives activation and stops doing so, when it stops receiving activation.

With MicroPsi node nets in the back, more complex experiments can now be thought of and implemented, that lead to more complex behaviours of the agent.

An issue that remains concerning is that the reaction time in between changes in the node net and the according changes of the agent's behaviour seems to be too long, for differing timing of the node net runner and the world runner.

This issue could be resolved, by making the bot faster—for example, by sending the visualisation image data directly to the web interface, without writing it to the hard disk first.

4.6 Summary

This chapter presented the implementation of the Minecraft interface for MicroPsi, which consists of the modules `minecraftClient`, `minecraftVisualisation` and `minecraftWorld`. The `minecraftClient` facilitates the communication with the Minecraft server; the `minecraftVisualisation` delivers OpenGL rendered image data of the environment and the `minecraftWorld` serves as the actual world adapter in between the `minecraftClient` and MicroPsi. Next, a case study has been performed and evaluated. A MicroPsi client had to move towards a block of a particular type. It showed that the interface is functional and ready for more sophisticated experiments. With the implementation of the interface and the following proof-of-concept case study the scope of this thesis has been completed. The last chapter contains a summary and an outlook towards future applications.



Figure 4.9: 1. Neither node net, nor agent show any activity, as the experiment did not start yet. 2. The sensors that indicate a diamond towards the negative x- and z-axes light up green—so do the according actuators. The agent starts moving. 3. The agent arrived close enough to the diamond, that the sensors stop forwarding activation. The agent stops moving.

Chapter 5

Conclusion

This thesis presented an implementation of an interface that makes it possible to use servers of the video game Minecraft as a simulation environment for the cognitive artificial intelligence framework MicroPsi 2. After several iterations and trying out different approaches and technologies, the interface is now functional. The proof-of-concept experiment with an agent seeking a particular block resulted in proofing that Minecraft is usable as a simulation environment.

In particular, an artificial player may now connect to a Minecraft world and be controlled by a MicroPsi node net. Because there is a wide variety of possible input data from Minecraft worlds and of possible actions artificial players can perform, corresponding data sources and targets may furthermore be defined freely inside the world adapter code.

The interface is in its core based on two open source software projects. The Python Minecraft bot framework **Spock** serves as the client that connects to Minecraft servers. The Python Minecraft clone **Minecraft** by Michael Fogleman serves as the foundation for the visualisation component.

The implemented case study shows that the interface as a whole works as expected. Because it so far serves foremost as a proof-of-concept, there is a lot of room left for improvement. The two main critical aspects are speed and robustness. Speeding the simulation up would require more profiling and finding bottle necks in the data flow. Most likely, at this time one of the major bottlenecks will be, that, for every iteration of the visualisation component, an image file is written to the hard drive.

The Minecraft protocol proved to be easily accessible, understandable and therefore gratifying to work with. It sets the hurdles for programming against Minecraft very low, so there might even be research projects with other scopes than AI simulation that could make use of Minecraft infrastructures.

Possible future applications of this project would be implementing data sources and targets for every aspect of a world and for every Minecraft action that one wants to experiment with. For example, these could include interacting with the environment by picking up objects, creating and destroying blocks, building, mining, interacting and fighting with other objects. Especially, in the future, multiple agents shall interact within the same environment and collaborate with each other.

List of Figures

2.1	Quads like this are the basic representational unit of the Psi theory (taken from [Bac09])	5
2.2	The basic architecture of action selection in Dörner's theory (taken from [Bac09])	6
2.3	The partonomic representation of a Cartoon face (taken from [Bac09]) .	6
2.4	Screenshot of Dörner's team's Delphi Pascal Psi implementation	8
2.5	Screenshots of the 3D visualisation component in MicroPsi (taken from [Bac09])	9
2.6	MicroPsi's simulation environment Berlin	10
2.7	The modular architecture of the MicroPsi framework makes it easy to extend. (taken from [Bac12])	11
2.8	The graphical editor is the primary interface to node nets. (taken from [Bac12])	12
3.1	A grass block (CC-BY-3.0 Mojang AB) [ima13]	17
3.2	A chunk (CC-BY-3.0 Mojang AB) [ima13]	17
3.4	Screenshots of Minecraft's basic game mechanics	18
3.3	The player (CC-BY-3.0 Mojang AB) [ima13]	18
3.5	An excerpt of Spock's data structures for packet-IDs and -structures .	21
3.6	Spock's function for decoding packets	21
3.7	A screenshot of Minecraft by Michael Fogleman	22
4.1	The new architecture of MicroPsi with the Minecraft interface	24
4.2	The <code>minecraftClient</code> 's event loop	25
4.3	By default, this function handles "Chunk Data (0x33) packets.	26
4.4	Checking for activation of the move-x data target and invoking the ac- cording action.	26
4.5	An overview of the most import classes of the <code>minecraftClient</code> in re- spect to the communication with a Minecraft Server	27
4.6	An image generated by the <code>minecraftVisualisation</code> component	28
4.7	Searching the current section for a diamond block and, if it is found, filling the concerning data sources with the distance towards it.	29
4.8	The node net set up for this experiment	29
4.9	1. Neither node net, nor agent show any activity, as the experiment did not start yet. 2. The sensors that indicate a diamond towards the negative x- and z-axes light up green—so do the according actuators. The agent starts moving. 3. The agent arrived close enough to the diamond, that the sensors stop forwarding activation. The agent stops moving.	31

List of Tables

3.1	The structure of the packet “Player Position (0x0B)” [pro13]	18
-----	--	----

Contents of the CD

The attached CD contains the following::

- this bachelor thesis in PDF format,
- the Python project MicroPsi including the Minecraft world adapter,
- the binaries of a Minecraft server that can be used to test the implementation.

Bibliography

- [Bac09] Joscha Bach. *Principles of Synthetic Intelligence PSI: An Architecture of Motivated Cognition*. Oxford University Press, Inc., New York, NY, USA, 1st edition, 2009.
- [Bac12] Joscha Bach. Micropsi 2: The next generation of the micropsi framework. In Joscha Bach, Ben Goertzel, and Matthew Iklé, editors, *AGI*, volume 7716 of *Lecture Notes in Computer Science*, pages 11–20. Springer, 2012.
- [BB] Gergo Balogh and Arpád Beszédes. Codemetropolis—a minecraft based collaboration tool for developers.
- [Dun11] Sean C. Duncan. Minecraft, beyond construction and survival. *Well Played*, 1(1):1–22, January 2011.
- [fog13a] August 2013.
- [fog13b] August 2013.
- [git13] August 2013.
- [GP07] Ben Goertzel and Cassio Pennachin. *Artificial general intelligence*. Springer, 2007.
- [ima13] August 2013.
- [mcf13] August 2013.
- [mcw13a] September 2013.
- [mcw13b] August 2013.
- [pro13] August 2013.
- [Psi13] September 2013.
- [pyg13] August 2013.
- [Rus09] Stuart Russell. Artificial intelligence: A modern approach author: Stuart russell, peter norvig, publisher: Prentice hall pa. 2009.
- [sky13] August 2013.